

S12 T01: Pipelines, grid search i text mining

```
In [78]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import category_encoders as ce
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.pipeline import make_pipeline
from sklearn.compose import make_column_transformer
from sklearn.compose import make_column_selector
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import Normalizer
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.utils import check_array
from scipy import sparse
from sklearn import datasets
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV
from sklearn.base import BaseEstimator
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import RandomizedSearchCV
import nltk
from nltk.tokenize import word_tokenize
from nltk.probability import FreqDist
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
import nltk
nltk.download('punkt')
nltk.download('stopwords')
from nltk import sent_tokenize
from textblob import TextBlob
```

```
[nltk_data] Downloading package punkt to /home/rusi/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /home/rusi/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

Exercici 1. Agafa el conjunt de dades que vulguis i realitza un pipeline i un gridsearch aplicant l'algorisme de Random Forest.

De l'Sprint07, carreguem les dades netes, sense nuls, amb l'històric de jugadors de la selecció espanyola de futbol absoluta masculina que han debutat (obtingudes a partir de la web bdfutbol.com). Recordem els noms de les columnes:

Sobrenom; Nom; Data Naixement; Lloc de Naixement; Província; País; Partits Jugats; Partits Titular; Partits
Complets; Partits Suplent; Partits Substituït; Partits Convocats (sense jugar); Partits Guanyats; Partits
Empetats; Partits Perduts; Minuts; Goles; Gols Penalt; Goles pròpia porta; Gols Encaixats; Targetes
grogues; Targetes vermelles; Edat inicial; Edat final; Alçada; Pes

```
In [2]: jugadores = pd.read_csv('///home/rusi/Escritorio/rubenIT/DataSources/jugadores00.csv')#imp
```

```
In [3]: #Imprimir les dades filtrades per comprovar la importació
print(jugadores.describe())
print(jugadores.head(10))
print(jugadores.tail(10))
```

	PJ	PT	PC	PS	PX	PG \
count	654.000000	654.000000	654.000000	654.000000	654.000000	654.000000
mean	14.155963	11.085627	8.006116	3.070336	3.056575	8.391437
std	22.460518	19.330256	14.271486	5.229901	7.115855	15.330149
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	2.000000	1.000000	1.000000	0.000000	0.000000	1.000000
50%	5.000000	4.000000	3.000000	1.000000	1.000000	3.000000
75%	16.000000	12.000000	9.000000	3.000000	3.000000	9.000000
max	180.000000	161.000000	125.000000	42.000000	59.000000	131.000000

	PE	PP	Min	G	GP \
count	654.000000	654.000000	654.000000	654.000000	654.000000
mean	3.333333	2.431193	1005.507645	1.960245	0.142202
std	4.831199	3.607972	1669.924268	5.165109	0.873092
min	0.000000	0.000000	1.000000	0.000000	0.000000
25%	0.000000	0.000000	90.000000	0.000000	0.000000
50%	1.000000	1.000000	360.000000	0.000000	0.000000
75%	4.000000	3.000000	1129.250000	1.000000	0.000000
max	33.000000	23.000000	13709.000000	59.000000	11.000000

	GPP	GE	TA	TR	EI	EF \
count	654.000000	654.000000	654.000000	654.000000	654.000000	654.000000
mean	0.019878	0.905199	0.917431	0.032110	23.949541	26.831804
std	0.139687	6.868723	2.419149	0.184904	2.782392	3.488660
min	0.000000	0.000000	0.000000	0.000000	17.000000	17.000000
25%	0.000000	0.000000	0.000000	0.000000	22.000000	25.000000
50%	0.000000	0.000000	0.000000	0.000000	24.000000	27.000000
75%	0.000000	0.000000	1.000000	0.000000	26.000000	29.000000
max	1.000000	100.000000	24.000000	2.000000	34.000000	36.000000

	Altura	Peso
count	654.000000	654.000000
mean	177.594801	73.915902
std	6.021862	5.713472
min	160.000000	60.000000
25%	173.000000	70.000000
50%	178.000000	74.000000
75%	181.750000	77.000000
max	197.000000	95.000000

	Apodo	Nombre	Fecha	Ciudad \
0	Marcos Vales	Marcos Vales Illanes	05/04/1975	A Coruña
1	Acuña	Juan Acuña Naya	13/02/1923	A Coruña
2	Martín	José María Martín Rodríguez	25/04/1924	A Coruña
3	Casilla	Francisco Casilla Cortés	02/10/1986	Alcover
4	Juan Sánchez	Juan Ginés Sánchez Romero	15/05/1972	Aldaia
5	Cucurella	Marc Cucurella Saseta	22/07/1998	Alella
6	Piquer	Vicente Piquer Mora	24/02/1935	Algar de Palancia
7	Ito	Antonio Álvarez Pérez	21/01/1975	Almendralejo
8	Planas II	Javier Planas Abad	03/07/1949	Almudévar
9	Josep Martínez	Josep Martínez Riera	27/05/1998	Alzira

	Provincia	País	PJ	PT	PC	PS	...	G	GP	GPP	GE	TA	TR	EI	EF	\
0	A Coruña	España	1	0	0	1	...	0	0	0	0	0	0	23	23	
1	A Coruña	España	1	0	0	1	...	0	0	0	1	0	0	18	18	
2	A Coruña	España	1	1	1	0	...	0	0	0	0	0	0	28	28	
3	Tarragona	España	1	0	0	1	...	0	0	0	1	0	0	28	28	
4	Valencia	España	1	0	0	1	...	0	0	0	0	0	0	26	26	
5	Barcelona	España	1	1	0	0	...	0	0	0	0	0	0	22	22	
6	Valencia	España	1	1	1	0	...	0	0	0	0	0	0	26	26	
7	Badajoz	España	1	0	0	1	...	0	0	0	0	0	0	23	23	
8	Huesca	España	1	1	1	0	...	0	0	0	0	1	0	25	25	
9	Valencia	España	1	0	0	1	...	0	0	0	0	0	0	23	23	

	Altura	Peso
0	181.0	77.0
1	179.0	88.0
2	176.0	74.0
3	192.0	83.0
4	173.0	72.0
5	172.0	68.0
6	173.0	71.0
7	175.0	70.0
8	174.0	74.0
9	191.0	78.0

[10 rows x 25 columns]

	Apodo	Nombre	Fecha	Ciudad	\
644	Fàbregas	Francesc Fàbregas Soler	04/05/1987	Arenys de Mar	
645	Fernando Torres	Fernando José Torres Sanz	20/03/1984	Fuenlabrada	
646	Xabi Alonso	Xabier Alonso Olano	25/11/1981	Tolosa	
647	Silva	David Josué Jiménez Silva	08/01/1986	Arguineguín	
648	Zubizarreta	Andoni Zubizarreta Urreta	23/10/1961	Vitoria-Gasteiz	
649	Iniesta	Andrés Iniesta Luján	11/05/1984	Fuentealbilla	
650	Busquets	Sergio Busquets Burgos	16/07/1988	Sabadell	
651	Xavi	Xavier Hernández Creus	25/01/1980	Terrassa	
652	Casillas	Iker Casillas Fernández	20/05/1981	Móstoles	
653	Sergio Ramos	Sergio Ramos García	30/03/1986	Camas	

	Provincia	País	PJ	PT	PC	PS	...	G	GP	GPP	GE	TA	TR	\
644	Barcelona	España	110	68	22	42	...	15	0	0	0	15	0	
645	Madrid	España	110	75	21	35	...	38	5	0	0	4	0	
646	Gipuzkoa	España	114	86	48	28	...	16	6	0	0	10	1	
647	Las Palmas	España	125	96	37	29	...	35	2	0	0	10	0	
648	Araba/Álava	España	126	125	106	1	...	0	0	1	100	2	1	
649	Albacete	España	131	105	47	26	...	13	1	0	0	4	0	
650	Barcelona	España	133	119	89	14	...	2	0	0	0	23	0	
651	Barcelona	España	133	108	64	25	...	13	0	0	0	9	0	
652	Madrid	España	167	154	125	13	...	0	0	0	93	2	0	
653	Sevilla	España	180	161	118	19	...	23	8	0	0	24	0	

	EI	EF	Altura	Peso
644	18	29	180.0	77.0
645	19	30	186.0	78.0
646	21	32	183.0	75.0
647	20	32	170.0	67.0
648	23	36	187.0	86.0
649	22	34	171.0	68.0
650	20	32	189.0	76.0
651	20	34	170.0	68.0
652	19	35	182.0	80.0
653	18	34	184.0	83.0

[10 rows x 25 columns]

Pipeline

1) Building a prototype Construïm el prototipus, millorant les dades del dataframe. Prescindirem dels features "Apodo", "Nombre", "Fecha" i "Ciudad", i convertirem en números "Provincia" i "País".

1.1) Encode the categorical variables.

```
In [4]: jugadores00=jugadores
```

```
In [5]: # create an object of the OneHotEncoder
OHE = ce.OneHotEncoder(cols=["Provincia","País"],use_cat_names=True)
# encode the categorical variables
jugadores00 = OHE.fit_transform(jugadores00)
```

```
In [6]: print(jugadores00.head())
print(jugadores00.info())
```

	Apodo	Nombre	Fecha	Ciudad	\
0	Marcos Vales	Marcos Vales Illanes	05/04/1975	A Coruña	
1	Acuña	Juan Acuña Naya	13/02/1923	A Coruña	
2	Martín	José María Martín Rodríguez	25/04/1924	A Coruña	
3	Casilla	Francisco Casilla Cortés	02/10/1986	Alcover	
4	Juan Sánchez	Juan Ginés Sánchez Romero	15/05/1972	Aldaia	

	Provincia_A Coruña	Provincia_Tarragona	Provincia_Valencia	\
0	1	0	0	
1	1	0	0	
2	1	0	0	
3	0	1	0	
4	0	0	1	

	Provincia_Barcelona	Provincia_Badajoz	Provincia_Huesca	...	G	GP	GPP	\
0	0	0	0	...	0	0	0	
1	0	0	0	...	0	0	0	
2	0	0	0	...	0	0	0	
3	0	0	0	...	0	0	0	
4	0	0	0	...	0	0	0	

	GE	TA	TR	EI	EF	Altura	Peso
0	0	0	0	23	23	181.0	77.0
1	1	0	0	18	18	179.0	88.0
2	0	0	0	28	28	176.0	74.0
3	1	0	0	28	28	192.0	83.0
4	0	0	0	26	26	173.0	72.0

```
[5 rows x 90 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 654 entries, 0 to 653
Data columns (total 90 columns):
```

#	Column	Non-Null Count	Dtype
0	Apodo	654 non-null	object
1	Nombre	654 non-null	object
2	Fecha	654 non-null	object
3	Ciudad	654 non-null	object
4	Provincia_A Coruña	654 non-null	int64
5	Provincia_Tarragona	654 non-null	int64
6	Provincia_Valencia	654 non-null	int64
7	Provincia_Barcelona	654 non-null	int64
8	Provincia_Badajoz	654 non-null	int64
9	Provincia_Huesca	654 non-null	int64
10	Provincia_Bizkaia	654 non-null	int64
11	Provincia_Málaga	654 non-null	int64
12	Provincia_La Rioja	654 non-null	int64
13	Provincia_Gipuzkoa	654 non-null	int64

14	Provincia_Extranjero	654	non-null	int64
15	Provincia_Burgos	654	non-null	int64
16	Provincia_Cádiz	654	non-null	int64
17	Provincia_Murcia	654	non-null	int64
18	Provincia_Zamora	654	non-null	int64
19	Provincia_Ceuta	654	non-null	int64
20	Provincia_Cáceres	654	non-null	int64
21	Provincia_Alicante	654	non-null	int64
22	Provincia_León	654	non-null	int64
23	Provincia_Segovia	654	non-null	int64
24	Provincia_Asturias	654	non-null	int64
25	Provincia_Granada	654	non-null	int64
26	Provincia_Palencia	654	non-null	int64
27	Provincia_Huelva	654	non-null	int64
28	Provincia_Sevilla	654	non-null	int64
29	Provincia_Jaén	654	non-null	int64
30	Provincia_Santa Cruz de Tenerife	654	non-null	int64
31	Provincia_Cantabria	654	non-null	int64
32	Provincia_Las Palmas	654	non-null	int64
33	Provincia_Lleida	654	non-null	int64
34	Provincia_Madrid	654	non-null	int64
35	Provincia_Toledo	654	non-null	int64
36	Provincia_Córdoba	654	non-null	int64
37	Provincia_Navarra	654	non-null	int64
38	Provincia_Lugo	654	non-null	int64
39	Provincia_Salamanca	654	non-null	int64
40	Provincia_Pontevedra	654	non-null	int64
41	Provincia_Valladolid	654	non-null	int64
42	Provincia_Araba/Álava	654	non-null	int64
43	Provincia_Zaragoza	654	non-null	int64
44	Provincia_Albacete	654	non-null	int64
45	Provincia_Almería	654	non-null	int64
46	Provincia_Castellón	654	non-null	int64
47	Provincia_Melilla	654	non-null	int64
48	Provincia_Girona	654	non-null	int64
49	Provincia_Ciudad Real	654	non-null	int64
50	Provincia_Teruel	654	non-null	int64
51	Provincia_Fernando Poo	654	non-null	int64
52	Provincia_Soria	654	non-null	int64
53	Provincia_Islas Baleares	654	non-null	int64
54	Provincia_Ourense	654	non-null	int64
55	Provincia_Ávila	654	non-null	int64
56	País_España	654	non-null	int64
57	País_Argentina	654	non-null	int64
58	País_Paraguay	654	non-null	int64
59	País_Suiza	654	non-null	int64
60	País_Italia	654	non-null	int64
61	País_Brasil	654	non-null	int64
62	País_Francia	654	non-null	int64
63	País_Dinamarca	654	non-null	int64
64	País_Guinea-Bisáu	654	non-null	int64
65	País_Hungría	654	non-null	int64
66	País_Guinea Ecuatorial	654	non-null	int64
67	País_Marruecos	654	non-null	int64
68	País_Alemania	654	non-null	int64
69	País_Mauritania	654	non-null	int64
70	País_Uruguay	654	non-null	int64
71	PJ	654	non-null	int64
72	PT	654	non-null	int64
73	PC	654	non-null	int64
74	PS	654	non-null	int64
75	PX	654	non-null	int64
76	PG	654	non-null	int64
77	PE	654	non-null	int64
78	PP	654	non-null	int64
79	Min	654	non-null	int64

```

80  G          654 non-null    int64
81  GP          654 non-null    int64
82  GPP         654 non-null    int64
83  GE          654 non-null    int64
84  TA          654 non-null    int64
85  TR          654 non-null    int64
86  EI          654 non-null    int64
87  EF          654 non-null    int64
88  Altura     654 non-null    float64
89  Peso       654 non-null    float64
dtypes: float64(2), int64(84), object(4)
memory usage: 460.0+ KB
None

```

El nou dataframe té 90 columnes, cosa que és massa nombrós pel nostre anàlisi. En comptes de fer-ho amb dummies, donarem valors numèrics a "Provincia" i "País" directament.

```

In [7]: number=LabelEncoder()
jugadors=jugadors.drop(["Apodo", "Nombre", "Fecha", "Ciudad"], axis=1)
jugadors["Provincia"]=number.fit_transform(jugadors["Provincia"].astype("str"))
jugadors["País"]=number.fit_transform(jugadors["País"].astype("str"))

```

```

In [8]: print(jugadors.iloc[0:50,:])
print(jugadors.info())

```

	Provincia	País	PJ	PT	PC	PS	PX	PG	PE	PP	...	G	GP	GPP	GE	TA	\
0	0	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
1	0	4	1	0	0	1	0	1	0	0	...	0	0	0	1	0	
2	0	4	1	1	1	0	0	1	0	0	...	0	0	0	0	0	
3	44	4	1	0	0	1	0	0	0	1	...	0	0	0	1	0	
4	47	4	1	0	0	1	0	0	1	0	...	0	0	0	0	0	
5	7	4	1	1	0	0	1	1	0	0	...	0	0	0	0	0	
6	47	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0	
7	6	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
8	23	4	1	1	1	0	0	1	0	0	...	0	0	0	0	1	
9	47	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
10	8	4	1	1	1	0	0	1	0	0	...	0	0	0	0	0	
11	34	4	1	0	0	1	0	0	0	1	...	0	0	0	0	0	
12	26	4	1	1	0	0	1	1	0	0	...	0	0	0	0	0	
13	26	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
14	8	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
15	19	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0	
16	19	4	1	0	0	1	0	0	0	1	...	0	0	0	0	0	
17	19	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
18	6	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0	
19	8	4	1	0	0	1	0	0	0	1	...	0	0	0	2	0	
20	8	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0	
21	7	4	1	0	0	1	0	1	0	0	...	1	0	0	0	0	
22	7	4	1	1	1	0	0	1	0	0	...	0	0	0	1	0	
23	7	4	1	1	1	0	0	1	0	0	...	0	0	0	0	0	
24	7	4	1	1	1	0	0	1	0	0	...	0	0	0	1	0	
25	7	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
26	7	4	1	0	0	1	0	0	1	0	...	0	0	0	0	0	
27	47	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
28	8	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
29	8	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
30	8	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0	
31	8	4	1	1	1	0	0	0	0	1	...	0	0	0	2	0	
32	17	1	1	1	1	0	0	1	0	0	...	0	0	0	0	0	
33	9	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
34	15	4	1	1	1	0	0	0	0	1	...	0	0	0	0	0	
35	26	4	1	0	0	1	0	0	0	1	...	0	0	0	0	0	
36	17	12	1	1	1	0	0	1	0	0	...	0	0	0	0	0	
37	33	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	
38	33	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0	

39	7	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0
40	49	4	1	1	0	0	1	0	0	1	...	0	0	0	0	0
41	47	4	1	1	0	0	1	0	1	0	...	0	0	0	0	0
42	12	4	1	1	1	0	0	1	0	0	...	0	0	0	0	0
43	34	4	1	1	0	0	1	0	0	1	...	0	0	0	0	0
44	14	4	1	1	1	0	0	0	0	1	...	0	0	0	4	0
45	8	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0
46	6	4	1	1	0	0	1	1	0	0	...	0	0	0	0	1
47	19	4	1	1	0	0	1	1	0	0	...	0	0	0	0	0
48	19	4	1	1	1	0	0	0	1	0	...	0	0	0	0	0
49	19	4	1	0	0	1	0	1	0	0	...	0	0	0	0	0

	TR	EI	EF	Altura	Peso
0	0	23	23	181.0	77.0
1	0	18	18	179.0	88.0
2	0	28	28	176.0	74.0
3	0	28	28	192.0	83.0
4	0	26	26	173.0	72.0
5	0	22	22	172.0	68.0
6	0	26	26	173.0	71.0
7	0	23	23	175.0	70.0
8	0	25	25	174.0	74.0
9	0	23	23	191.0	78.0
10	0	27	27	186.0	80.0
11	0	20	20	172.0	72.0
12	0	23	23	186.0	80.0
13	0	23	23	179.0	74.0
14	0	25	25	177.0	70.0
15	0	29	29	173.0	71.0
16	0	27	27	185.0	84.0
17	0	29	29	177.0	75.0
18	0	29	29	174.0	72.0
19	0	24	24	178.0	76.0
20	0	25	25	170.0	66.0
21	0	23	23	177.0	69.0
22	0	26	26	183.0	80.0
23	0	22	22	178.0	66.0
24	0	33	33	180.0	80.0
25	0	23	23	168.0	62.0
26	0	20	20	170.0	70.0
27	0	23	23	172.0	71.0
28	0	26	26	183.0	83.0
29	0	21	21	186.0	79.0
30	0	25	25	176.0	73.0
31	0	27	27	180.0	77.0
32	0	27	27	178.0	72.0
33	0	21	21	183.0	81.0
34	0	23	23	184.0	81.0
35	0	21	21	180.0	70.0
36	0	25	25	168.0	67.0
37	0	28	28	181.0	73.0
38	0	23	23	197.0	90.0
39	0	22	22	172.0	68.0
40	0	25	25	181.0	80.0
41	0	22	22	182.0	78.0
42	0	29	29	182.0	74.0
43	0	21	21	172.0	69.0
44	0	28	28	185.0	82.0
45	0	25	25	176.0	70.0
46	0	21	21	176.0	71.0
47	0	22	22	180.0	74.0
48	0	21	21	177.0	75.0
49	0	28	28	182.0	78.0

[50 rows x 21 columns]
<class 'pandas.core.frame.DataFrame'>

```

RangeIndex: 654 entries, 0 to 653
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Provincia   654 non-null    int64
1   País        654 non-null    int64
2   PJ          654 non-null    int64
3   PT          654 non-null    int64
4   PC          654 non-null    int64
5   PS          654 non-null    int64
6   PX          654 non-null    int64
7   PG          654 non-null    int64
8   PE          654 non-null    int64
9   PP          654 non-null    int64
10  Min         654 non-null    int64
11  G           654 non-null    int64
12  GP          654 non-null    int64
13  GPP         654 non-null    int64
14  GE          654 non-null    int64
15  TA          654 non-null    int64
16  TR          654 non-null    int64
17  EI          654 non-null    int64
18  EF          654 non-null    int64
19  Altura      654 non-null    float64
20  Peso        654 non-null    float64
dtypes: float64(2), int64(19)
memory usage: 107.4 KB
None

```

1.2) Scale data

Podem escalar totes les columnes (entre 0 i 1), excepte "Peso", que el normalitzem (mitjana=0 i desviació=1). "Altura" és el target i no el tractem.

```

In [9]: #Estandarització i eliminació target "Altura"
jugadors01 = jugadors.drop(["Provincia", "País", "Peso", "Altura"], axis=1)
jugadors01_norm = (jugadors01-jugadors01.min())/(jugadors01.max()-jugadors01.min())
#jugadors01_norm = jugadors.drop(["Altura"], axis=1)

```

```

In [10]: print(jugadors01_norm.head())
print(jugadors01_norm.info())

```

```

      PJ      PT      PC      PS      PX      PG      PE      PP      Min \
0  0.0  0.000000  0.000  0.02381  0.0  0.007634  0.000000  0.000000  0.000000
1  0.0  0.000000  0.000  0.02381  0.0  0.007634  0.000000  0.000000  0.001094
2  0.0  0.006211  0.008  0.00000  0.0  0.007634  0.000000  0.000000  0.006493
3  0.0  0.000000  0.000  0.02381  0.0  0.000000  0.000000  0.043478  0.000875
4  0.0  0.000000  0.000  0.02381  0.0  0.000000  0.030303  0.000000  0.000802

```

```

      G      GP      GPP      GE      TA      TR      EI      EF
0  0.0  0.0  0.0  0.00  0.0  0.0  0.352941  0.315789
1  0.0  0.0  0.0  0.01  0.0  0.0  0.058824  0.052632
2  0.0  0.0  0.0  0.00  0.0  0.0  0.647059  0.578947
3  0.0  0.0  0.0  0.01  0.0  0.0  0.647059  0.578947
4  0.0  0.0  0.0  0.00  0.0  0.0  0.529412  0.473684

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 654 entries, 0 to 653
```

```
Data columns (total 17 columns):
```

```

#   Column      Non-Null Count  Dtype
---  -
0   PJ          654 non-null    float64
1   PT          654 non-null    float64
2   PC          654 non-null    float64
3   PS          654 non-null    float64

```



```

4   PX      654 non-null    float64
5   PG      654 non-null    float64
6   PE      654 non-null    float64
7   PP      654 non-null    float64
8   Min     654 non-null    float64
9   G       654 non-null    float64
10  GP      654 non-null    float64
11  GPP     654 non-null    float64
12  GE      654 non-null    float64
13  TA      654 non-null    float64
14  TR      654 non-null    float64
15  EI      654 non-null    float64
16  EF      654 non-null    float64

```

```

dtypes: float64(17)
memory usage: 87.0 KB
None

```

```

In [11]: #Normalització
jugadors02=jugadors.loc[:,["Provincia","País","Peso"]]
ss = StandardScaler()
jugadors03 = ss.fit_transform(jugadors02.to_numpy())
jugadors03 = pd.DataFrame(jugadors03, columns=["Provincia","País","Peso"])

```

```

In [12]: print(jugadors03.head())
print(jugadors03.info())

```

```

      Provincia      País      Peso
0  -1.524230 -0.070183  0.540207
1  -1.524230 -0.070183  2.466955
2  -1.524230 -0.070183  0.014730
3   1.476060 -0.070183  1.591161
4   1.680626 -0.070183 -0.335587
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 654 entries, 0 to 653
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Provincia  654 non-null    float64
1   País       654 non-null    float64
2   Peso       654 non-null    float64
dtypes: float64(3)
memory usage: 15.5 KB
None

```

```

In [13]: jugadors04 = pd.concat((jugadors01_norm,jugadors03.loc[:,["Provincia","País","Peso"]]),
jugadors04= pd.concat((jugadors04,jugadors.loc[:, "Altura"]), 1)

```

```

/tmp/ipykernel_42366/669182810.py:1: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
jugadors04 = pd.concat((jugadors01_norm,jugadors03.loc[:,["Provincia","País","Peso"]]), 1)
/tmp/ipykernel_42366/669182810.py:2: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only
jugadors04= pd.concat((jugadors04,jugadors.loc[:, "Altura"]), 1)

```

```

In [14]: print(jugadors04.head())
print(jugadors04.info())

```

```

      PJ      PT      PC      PS      PX      PG      PE      PP      Min \
0  0.0  0.000000  0.000  0.02381  0.0  0.007634  0.000000  0.000000  0.000000
1  0.0  0.000000  0.000  0.02381  0.0  0.007634  0.000000  0.000000  0.001094
2  0.0  0.006211  0.008  0.00000  0.0  0.007634  0.000000  0.000000  0.006493
3  0.0  0.000000  0.000  0.02381  0.0  0.000000  0.000000  0.043478  0.000875
4  0.0  0.000000  0.000  0.02381  0.0  0.000000  0.030303  0.000000  0.000802

      G  ...  GPP      GE  TA  TR      EI      EF  Provincia      País \

```

```

0  0.0  ...  0.0  0.00  0.0  0.0  0.352941  0.315789  -1.524230  -0.070183
1  0.0  ...  0.0  0.01  0.0  0.0  0.058824  0.052632  -1.524230  -0.070183
2  0.0  ...  0.0  0.00  0.0  0.0  0.647059  0.578947  -1.524230  -0.070183
3  0.0  ...  0.0  0.01  0.0  0.0  0.647059  0.578947   1.476060  -0.070183
4  0.0  ...  0.0  0.00  0.0  0.0  0.529412  0.473684   1.680626  -0.070183

```

```

      Peso  Altura
0  0.540207  181.0
1  2.466955  179.0
2  0.014730  176.0
3  1.591161  192.0
4 -0.335587  173.0

```

```

[5 rows x 21 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 654 entries, 0 to 653
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PJ          654 non-null    float64
1   PT          654 non-null    float64
2   PC          654 non-null    float64
3   PS          654 non-null    float64
4   PX          654 non-null    float64
5   PG          654 non-null    float64
6   PE          654 non-null    float64
7   PP          654 non-null    float64
8   Min         654 non-null    float64
9   G           654 non-null    float64
10  GP          654 non-null    float64
11  GPP         654 non-null    float64
12  GE          654 non-null    float64
13  TA          654 non-null    float64
14  TR          654 non-null    float64
15  EI          654 non-null    float64
16  EF          654 non-null    float64
17  Provincia   654 non-null    float64
18  País        654 non-null    float64
19  Peso        654 non-null    float64
20  Altura      654 non-null    float64
dtypes: float64(21)
memory usage: 107.4 KB
None

```

1.3) Model building

Fem servir el model predictiu de regressió Random Forest, i trobem el RMSE (Root Mean Squared Error).

```

In [15]: X = jugadors04.drop(columns=["Altura"])
        y = jugadors04["Altura"]

        # randomly split the data
        X, test_X, y, test_y = train_test_split(X,y,test_size=0.3,random_state=42)

        # shape of train and test splits
        X.shape, test_X.shape, y.shape, test_y.shape

```

```

Out[15]: ((457, 20), (197, 20), (457,), (197,))

```

```

In [16]: # create an object of the RandomForestRegressor
        model_RFR = RandomForestRegressor(max_depth=10)

        # fit the model with the training data
        model_RFR.fit(X, y)

```

```
# predict the target on train and test data
predict_train = model_RFR.predict(X)
predict_test = model_RFR.predict(test_X)

# Root Mean Squared Error on train and test data
print('RMSE on train data: ', mean_squared_error(y, predict_train)**(0.5))
print('RMSE on test data: ', mean_squared_error(test_y, predict_test)**(0.5))
```


RMSE on train data: 1.6242631713439897

RMSE on test data: 3.6081435794826793

1.4) Feature Importance

Veiem quina és la importància de cada un dels atributs per predir el target "Altura".

```
In [17]: plt.figure(figsize=(10,7))
feat_importances = pd.Series(model_RFR.feature_importances_, index = X.columns)
feat_importances.nlargest(20).plot(kind='barh');
```



Amb 6 features arribaríem a prop del 80% del total d'importància. Fixem aquesta dada i comparem el RMSE per contrastar la millora amb aquesta reducció.

```
In [18]: X = jugadors04.loc[:,["Peso", "Provincia", "PS", "EF", "EI", "Min"]]
y = jugadors04["Altura"]

# randomly split the data
X, test_X, y, test_y = train_test_split(X, y, test_size=0.3, random_state=42)

# shape of train and test splits
X.shape, test_X.shape, y.shape, test_y.shape
```

```
Out[18]: ((457, 6), (197, 6), (457,), (197,))
```

```
In [19]: # create an object of the RandomForestRegressor
model_RFR = RandomForestRegressor(max_depth=10)
```

```
# fit the model with the training data
model_RFR.fit(X, y)

# predict the target on train and test data
predict_train = model_RFR.predict(X)
predict_test = model_RFR.predict(test_X)

# Root Mean Squared Error on train and test data
print('RMSE on train data: ', mean_squared_error(y, predict_train)**(0.5))
print('RMSE on test data: ', mean_squared_error(test_y, predict_test)**(0.5))
```

RMSE on train data: 1.660670329283194

RMSE on test data: 3.690543837014724

El model amb la reducció de **features** puja molt lleugerament el RMSE, així que aprofitem el nou dataframe per aplicar-ho en el pipeline (1.624 vs 1.660 en el train, i 3.608 vs 3.690 en el test).

1.5) Identify features to build the Machine Learning pipeline

```
In [20]: jugadors05=jugadors04.loc[:,["Altura","Peso","Provincia","PS","EF","EI","Min"]]
```

```
In [21]: print(jugadors05.head())
print(jugadors05.info())
```

```
      Altura      Peso  Provincia      PS      EF      EI      Min
0   181.0    0.540207  -1.524230  0.02381  0.315789  0.352941  0.000000
1   179.0    2.466955  -1.524230  0.02381  0.052632  0.058824  0.001094
2   176.0    0.014730  -1.524230  0.000000  0.578947  0.647059  0.006493
3   192.0    1.591161   1.476060  0.02381  0.578947  0.647059  0.000875
4   173.0   -0.335587   1.680626  0.02381  0.473684  0.529412  0.000802
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 654 entries, 0 to 653
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Altura      654 non-null    float64
1   Peso        654 non-null    float64
2   Provincia    654 non-null    float64
3   PS          654 non-null    float64
4   EF          654 non-null    float64
5   EI          654 non-null    float64
6   Min         654 non-null    float64
dtypes: float64(7)
memory usage: 35.9 KB
None
```

2) Pipeline design

Hem identificat els següents passos de preprocessat per crear el nostre model de pipeline: 1)Drop columns. 2)Transform column (labelencoder). 3)Scale data. 4)Normalització.

```
In [22]: jugadors = pd.read_csv('//home/rusi/Escritorio/rubenIT/DataSources/jugadores00.csv')#imp
```

```
In [23]: # separate the independent and target variables
X02 = jugadors.drop(columns=["Altura"])
y02 = jugadors["Altura"]
```

```
In [24]: train_x, test_x, train_y, test_y = train_test_split(X02, y02, test_size=0.3, random_state=
```

```
In [25]: numeric_preprocessor = Pipeline(
    steps=[
        ("imputation_mean", SimpleImputer(missing_values=np.nan, strategy="mean")),
```



```
['PS', 'Min', 'EI', 'EF']]])),  
('randomforestregressor',  
 RandomForestRegressor(max_depth=10, random_state=42))])])
```

```
In [26]: # fit the pipeline with the training data  
pipe.fit(train_x, train_y)  
# predict target values on the training data  
pipe.predict(test_x)
```

```
Out[26]: array([175.70934259, 182.62275806, 184.61037535, 180.46492238,  
170.18388473, 181.04930736, 175.30013095, 170.73893078,  
170.73826255, 178.22134329, 172.7170088 , 172.96667349,  
176.73625134, 182.16222018, 176.70949405, 172.55355411,  
171.28748551, 175.18344977, 181.37762616, 171.78751732,  
183.39931668, 181.16247835, 181.79668817, 181.33066089,  
174.56070859, 175.37616553, 178.7142957 , 185.27348798,  
182.16330118, 181.04783697, 181.34951104, 179.71060396,  
173.97952544, 180.85123427, 178.48788979, 181.5993148 ,  
183.22544473, 174.74031947, 170.66501242, 178.46949287,  
169.65351367, 170.53061029, 172.018475 , 188.76316767,  
177.00195735, 180.57387038, 181.47389081, 178.38415521,  
174.96053548, 177.10657937, 176.16495238, 171.65366041,  
171.07559207, 171.02685326, 182.18522598, 169.45631047,  
182.65871301, 170.53725973, 175.76928863, 183.95496336,  
175.83548135, 177.96691955, 174.41996766, 171.1766132 ,  
169.06847222, 181.77740227, 169.62248369, 187.90526399,  
180.48785562, 181.20421831, 176.46749331, 175.22691993,  
174.07071871, 175.32938395, 177.43276428, 184.34861785,  
174.17520599, 178.51552404, 175.07894689, 183.4697094 ,  
177.06508308, 171.87961166, 182.1472239 , 180.16039857,  
171.16919851, 185.13050979, 169.47788166, 182.4696656 ,  
180.52105256, 176.87118235, 188.29472078, 169.95846078,  
178.33969122, 183.72580188, 176.5355347 , 184.51028431,  
187.93351282, 184.86294514, 170.35129575, 184.85678236,  
171.06296413, 180.16670976, 176.9005173 , 181.69344763,  
179.62538082, 181.94616793, 184.38561622, 186.23125857,  
180.04229917, 169.95613297, 173.65130767, 184.33304477,  
176.7964599 , 174.86317835, 181.01827842, 184.57695858,  
178.6910504 , 177.94267544, 179.50632988, 188.33526399,  
170.88980423, 177.79110949, 182.66744012, 179.811 ,  
175.14726626, 175.75420971, 180.65572222, 181.03099767,  
174.74295569, 179.61448914, 183.37850509, 183.10208857,  
178.4818889 , 175.01620144, 182.45703965, 187.29387191,  
175.26125397, 179.82677778, 180.80844372, 176.09262165,  
170.21541882, 182.08104202, 179.69972449, 174.11021197,  
180.07232937, 177.63708112, 171.62038814, 180.65467322,  
171.92103598, 182.74038885, 175.29847646, 173.05831802,  
172.09100553, 182.4357135 , 176.43702194, 176.17240629,  
179.41633176, 171.70437288, 176.08597554, 170.78820497,  
174.00777177, 179.37016667, 173.28406917, 180.66614059,  
176.45800619, 175.4774509 , 183.88111812, 170.38698255,  
171.23930135, 177.2563169 , 187.33748252, 170.60183886,  
178.3435 , 176.89428005, 185.81909363, 176.94138354,  
182.39317663, 181.18210509, 175.06751551, 186.44544239,  
171.2681509 , 182.44814201, 171.56614286, 169.83255565,  
176.48228878, 182.57381867, 184.76933874, 176.03640042,  
178.59716247, 170.2688297 , 169.0827544 , 182.05125032,  
176.56255556, 176.29161625, 183.97259376, 178.70810726,  
182.00939749])
```

```
In [27]: # read the test data  
test_data = jugadores = pd.read_csv('///home/rusi/Escritorio/rubenIT/DataSources/jugadores  
predict_test = pipe.predict(test_x)
```

```
In [35]: print("Mitja train_y: ", np.mean(train_y))
```

```
print("Mitja test_y: ", np.mean(test_y))
print("Mitja predict_train: ", np.mean(predict_train))
print("Mitja predict_test: ", np.mean(predict_test))
print("Mitja Altura jugadors: ", np.mean(jugadors.Altura))
```

```
Mitja train_y: 177.50765864332604
Mitja test_y: 177.7969543147208
Mitja predict_train: 177.5076879073611
Mitja predict_test: 177.81352830461853
Mitja Altura jugadors: 177.5948012232416
```

```
In [36]: print('RMSE on train data: ', mean_squared_error(train_y, predict_train)**(0.5))
print('RMSE on test data: ', mean_squared_error(test_y, predict_test)**(0.5))
```

```
RMSE on train data: 2.7631782507651415
RMSE on test data: 3.654590587410196
```

Com podem veure, el valor "Altura" s'apropa molt a la mitjana dels valors d'entrenament i de tota la sèrie, 177.507 vs 177.598 vs 177.595. La predicció amb un 30% de les dades per fer el testeig, dona un valor de 177.853, molt a prop dels valors mitjans. Podem afirmar que el resultat obtingut és força bo en aquest sentit.

En quant a l'error RMSE, incrementa notablement en el train (2.763 vs 1.660), i és molt semblant en el test (3.654 vs 3.690). Té la seva lògica, ja que a mesura que hi ha més dades, tindrem més dispersió que fomentarà un increment en el RMSE.

Grid search

Grid Search Parameter Tuning. Grid search is an approach to parameter tuning that will methodically build and evaluate a model for each combination of algorithm parameters specified in a grid.

```
In [40]: # Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the random grid
random_grid = {'randomforestregressor__n_estimators': n_estimators,
               'randomforestregressor__max_features': max_features,
               'randomforestregressor__max_depth': max_depth,
               'randomforestregressor__min_samples_split': min_samples_split,
               'randomforestregressor__min_samples_leaf': min_samples_leaf,
               'randomforestregressor__bootstrap': bootstrap}
```

```
In [41]: # Use the random grid to search for best hyperparameters
# First create the base model to tune
rf = RandomForestRegressor()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = pipe, param_distributions = random_grid,
                               n_iter = 30, cv = 3, random_state=42, n_jobs = -1)
```



```
1400,  
1600,  
1800,  
2000]],
```

```
random_state=42)
```

```
In [43]: rf_random.best_params_
```

```
Out[43]: {'randomforestregressor__n_estimators': 1800,  
         'randomforestregressor__min_samples_split': 10,  
         'randomforestregressor__min_samples_leaf': 4,  
         'randomforestregressor__max_features': 'auto',  
         'randomforestregressor__max_depth': 90,  
         'randomforestregressor__bootstrap': True}
```

```
In [44]: # Best estimator  
best_random = rf_random.best_estimator_  
  
#Predictions  
predict_train = best_random.predict(train_x)  
predict_test = best_random.predict(test_x)  
  
# Root Mean Squared Error on train data  
print('RMSE on train data: ', mean_squared_error(train_y, predict_train)**(0.5))  
print('RMSE on test data: ', mean_squared_error(test_y, predict_test)**(0.5))
```

```
RMSE on train data:  2.7631782507651415  
RMSE on test data:  3.654590587410196
```

L'error RMSE es manté en el mateix valor. Això podria ser degut a què les mostres són escasses, i que no té sentit aplicar una millora d'aquest tipus.

Exercici 2. Agafa un text en anglès que vulguis, i calcula'n la freqüència de les paraules


De la següent pàgina web, descarreguem un capítol del llibre de Jane Austin "Sense and Sensibility":

<https://www.fulltextarchive.com/>

```
In [45]: # Llegir llibre  
book_file = open('/home/rusi/Escritorio/rubenIT/DataSources/sense_and_sensibility.txt',  
book = book_file.read())
```

```
In [65]: # Word freq  
tokenized_book = word_tokenize(book)  
# Remove punctuation  
tokenized_book = [word for word in tokenized_book if word.isalnum()]  
# Freq  
fdist = FreqDist(tokenized_book)  
# 30 most common words  
fig, ax1 = plt.subplots(figsize = (15, 10))  
fdist.plot(30, cumulative=False,  
          title="30 most often words in Jane Austin's -- Sense and Sensibility")  
plt.show()
```

30 most often words in Jane Austin's -- Sense and Sensibility



Exercici 3. Treu les stopwords i realitza stemming al teu conjunt de dades.


```
In [75]: # Stop words list
stop_words=set(stopwords.words("english"))

# Filter
filtered_book=[]
for w in tokenized_book:
    if w not in stop_words:
        filtered_book.append(w)

# Word freq
fdist = FreqDist(filtered_book)

# 30 most common words
fig, ax1 = plt.subplots(figsize = (15, 10))
fdist.plot(30,cumulative=False,title="Top 30 stop words in Jane Austin's -- Sense and Se
plt.show()
```

Top 30 stop words in Jane Austin's -- Sense and Sensibility




```
In [74]: # Stemming
ps = PorterStemmer()

stemmed_book=[]
for w in filtered_book:
    stemmed_book.append(ps.stem(w))

# Word freq
fdist = FreqDist(stemmed_book)

# 30 most common words
fig, ax1 = plt.subplots(figsize = (15, 10))
fdist.plot(30,cumulative=False,title="Top 30 stemmed words in Jane Austin's -- Sense and
plt.show()
```

Top 30 stemmed words in Jane Austin's -- Sense and Sensibility



Exercici 3. Realitza sentiment analysis al teu conjunt de dades

Sentiment Analysis

The *sentiment* property returns a *namedtuple* of the form *Sentiment(polarity, subjectivity)*. The *polarity* score is a float within the range $[-1.0, 1.0]$. The *subjectivity* is a float within the range $[0.0, 1.0]$ where 0.0 is very objective and 1.0 is very subjective.

```
In [79]: # Polarity of the text
book_sent = TextBlob(book)
book_sent.sentiment
```

```
Out[79]: Sentiment(polarity=0.15604552058239224, subjectivity=0.535295224185392)
```

Com es tracta d'una novel·la, la quantitat de paraules emprades són molt variades, i no demostren una inclinació cap a una determinada sensació o sentiment. Per aquest motiu els paràmetres de neutralitat i subjectivitat es troben en un punt intermig.

```
In [ ]:
```